

DO CONTEXT ENGINES HELP?

RETRIEVAL QUALITY, DETERMINISM, AND AGENT UTILITY

UNDER MATCHED BASELINES AND AUDITED EXPOSURE

Anonymous authors

Paper under double-blind review

ABSTRACT

Coding agents are increasingly wrapped in “context engines” that index a repository and return the files the model might need. We evaluate one open-source engine, Delphi, against lexical baselines, a matched ladder of conventional retrievers, and two hosted engines, under a protocol that records at the case level what each evaluation set could have influenced. Three findings temper our own earlier claims. *Exposure*: the largest partition we had called “untouched”—220 Agent Retrieval Bench cases—was re-drawn from a pool whose cases had all been scored in an earlier round and whose aggregate results gated which retrieval components shipped; we relabel it validation evidence and rest confirmatory claims on sets never scored before their single run: 62 SWE-bench Verified localization instances, 18 repository-disjoint commit-to-files cases, and a further 100 Verified instances drawn under a fresh salt. *Baselines*: Delphi beats BM25 and the official lexical ranker on every set, but a conventional dense+BM25 hybrid built from the same embedding model, with Delphi’s own cross-encoder, listwise reranker, and query expansion attached, matches or exceeds it; on the 18-case independent final the hybrid’s Recall@20 lead (+0.120) and the full conventional rung’s MRR lead (+0.111) have repository-cluster intervals excluding zero. on the 62-case SWE-bench set Delphi leads that stack by +0.083 MRR with an interval through zero [−0.107, +0.203], and the reranking head alone accounts for most of its margin over lexical systems. What is specific to Delphi—structural candidate branches, tuned weights, quoted-path demotion—does not separate from that conventional stack. *Contracts*: routing every engine’s retrieved evidence, including the hosted synthesis engine’s own retrieved sources, through one frozen synthesis stage yields a four-way documentation tie (0.625/0.617/0.583/0.575) sitting 0.08–0.13 above a 0.492 no-retrieval floor; no engine-versus-engine interval excludes zero. On like-for-like retrieved-set stability the hosted synthesis engine is fully stable and Delphi is 98% stable; the “0%” we had reported for it measured generated text. A preregistered executable pilot on the same 62 SWE-bench instances (mini-SWE-agent, fixed budget) finds that a Delphi seed raises verified repairs from 25 to 31 (+0.097 [+0.016, +0.194]), but a random-file control rises to 30, so the value of naming the right files is not separable from the value of the prompt block at this size; the conventional ladder’s seed, despite similar gold coverage, resolved 23. We claim no state of the art. The contribution is the protocol, the negative and null results, and an exposure ledger that lets a reader check which number can bear which claim.

1 INTRODUCTION

The pitch for a context engine is simple. Index the repo. Embed the chunks. When the agent asks a question, return the files that matter. The pitch for *not* building one is simpler: `rg` is deterministic,

takes 30 milliseconds, and already found `src/_pytest/assertion/rewrite.py` when we asked where `pytest` rewrites asserts.

We built Delphi as a local-first retrieval stack—hybrid lexical/vector search, query expansion, a cross-encoder, a listwise rerank—and spent two evaluation rounds measuring it. The first round reported large wins over lexical baselines on a partition we called untouched. External review asked three questions we could not answer from the draft: what information from the earlier round was available while the shipped system was chosen; how much of the margin survives against a conventional dense-plus-lexical hybrid that uses the same reranker; and whether any of it changes what a coding agent actually repairs. This paper is the answer, and much of it is unflattering.

1. **RQ1 (utility).** Holding the agent, its tools, and its budget fixed, does a retrieval seed change what it submits—and does it change verified repair?
2. **RQ2 (ranking).** Against lexical baselines, against a matched conventional hybrid, and against hosted engines, where does Delphi lead, tie, or trail, on cases that no development decision could see?
3. **RQ3 (determinism).** Where does run-to-run variance come from, and when the same quantity is measured for every engine, who is actually stable?

Short answers. **RQ2:** Delphi leads lexical baselines everywhere; against the conventional ladder it does not lead on the unused sets, and on the independent final it trails with intervals excluding zero (§4). **RQ3:** the variance was ours and is fixed in-process, but the like-for-like comparison shows the hosted synthesis engine’s retrieval is at least as stable as ours (§5). **RQ1:** seeds move a closed-tool agent’s file selection, replicating the benchmark’s own intervention; on executable repair a Delphi seed helps a small agent at $n=62$, a random-file block helps almost as much, and the conventional seed does not help (§7).

2 RELATED WORK

SWE-bench established executable repository-level issue resolution (Jimenez et al., 2024), while SWE-agent and Agentless expose repository navigation and localization as explicit stages (Yang et al., 2024; Xia et al., 2024). LocAgent studies a graph-guided localization agent and reports that better localization can improve multi-attempt repair (Chen et al., 2025). Repository-level code completion provides complementary evidence that cross-file retrieval matters once the edit location is known: RepoCoder alternates retrieval and generation (Zhang et al., 2023); CodeRAG adds multi-path retrieval and preference-aligned reranking (Zhang et al., 2025). Dense retrieval, late interaction, and hypothetical-document expansion motivate the components in practical context engines (Karpukhin et al., 2020; Khattab & Zaharia, 2020; Gao et al., 2023), while BM25 remains a strong exact-term baseline for code (Robertson & Zaragoza, 2009).

The closest work measures context acquisition directly. ContextBench provides human-verified file, block, and line contexts for issue-resolution tasks and scores agent trajectories (Li et al., 2026). CORE-Bench scales requirement-driven retrieval across code understanding, edit localization, and broader context (Zhang et al., 2026a). SWE-Explore derives audited line-level targets from successful trajectories (Zhang et al., 2026b). Agent Retrieval Bench (ARB) contributes workflow-derived next-context targets, base-commit hygiene, budget-aware file metrics, open-embedding and RepoMap baselines, and a closed-tool seed intervention with no-seed, random, retrieval-derived, and oracle arms (Qin & Xie, 2026). Our closed-tool experiment (§7) is a replication of that intervention with an additional seed, not a new design. What we add to this line is narrower: an exposure ledger that classifies every case by what could have influenced it, a component-matched baseline ladder, matched output contracts for hosted engines, like-for-like determinism, and a preregistered executable pilot on the same tasks as the ranking evaluation.

3 PROTOCOL

Exposure classes. Every evaluation set is assigned a class in a case-level ledger (Appendix B). **C0:** never scored by any system before its single confirmatory run, drawn from a pool no development decision could see. **C1:** development. **C2:** re-partitioned from a pool whose cases were all scored

in an earlier round, where that round’s aggregates influenced the shipped system. The ARB round-3 “final” is C2: all 427 ARB v2 cases were scored in July 2026; the July held-out aggregates were the acceptance test for query expansion, listwise reranking, the cross-encoder blend, reciprocal-rank fusion, and the path-affinity branch, and the July per-workflow rates motivated the quoted-path demotion fix. Round 3 then drew a new 25/75 split from the same pool, excluded 50 legacy-exposed IDs, and queried the 220-case partition once. Re-drawing does not un-expose a case. Round-2 per-case outputs are not retained, so per-case tuning can neither be ruled in nor out from the archive; we therefore treat the partition as validation evidence throughout and make no confirmatory claim on it. The C0 sets are: 62 stratified SWE-bench Verified localization instances (12 repositories; patch-marker leakage audit 0/62); an 18-case commit-to-files final over a repository-disjoint corpus locked before any scoring; and 100 further Verified instances drawn on 2026-09-09 under a fresh salt from the 438 instances no round had used. A 40-case DS-1000 development subset drives the documentation track and is C1. The ledger controls author exposure; it cannot control whether the foundation models inside any engine saw these public repositories during training, a caveat shared by every arm.

Queries, corpus, metrics. Queries are the official ARB gold-free structured objects, serialized with sorted keys. File text comes from bare git clones at each case’s `base_commit` (materialized snapshots; binary and non-UTF-8 blobs skipped). Metrics are the official ARB `sample_metrics` (MRR, Recall@ k) plus BCY@8k, the budget-constrained yield from packing ranked files into an 8,000-token context. Documentation cases score *identifier hit*: whether the returned text contains the case’s gold API identifier.

Freeze and attestation. All Delphi numbers come from one frozen commit (91d76c1) with exact vector scan, API top- $k=20$, and response-level serving-configuration attestation: every response carries the retrieval configuration actually served, and a run aborts on the first mismatch with its declared configuration. The frozen configuration is reproduced in full in Appendix A. Paired deltas carry cluster bootstrap 95% intervals (20,000 resamples, seed 1042; repository clusters for repository tracks, case clusters for documentation), plus exact McNemar tests on any-gold@20 movement.

Matched baseline ladder. The lexical comparators (whole-file BM25, ARB’s official lexical ranker, and their development-tuned fusion) cannot say how much of a margin is Delphi’s engineering and how much is a learned reranker on top of anything. We therefore add four conventional systems, each adding one Delphi component to the previous one, over the official ARB chunks (file chunk plus symbol chunks) of the identical corpus: **dense**, `text-embedding-3-small`—the embedding model inside Delphi—max-pooled to files; **hybrid**, reciprocal-rank fusion ($k=60$, equal weights, untuned) of the dense ranking and BM25; **+ rerankers**, the hybrid’s top 50 narrowed to 30 through Delphi’s own cross-encoder (`ms-marco-MiniLM-L-6-v2`, blend 0.4) and then Delphi’s own listwise `gpt-4o` reranker over the top 20 (seed 1042, identical prompt); **+ expansion**, the same with Delphi’s hypothetical-document expansion (`gpt-4o-mini`, identical prompt and gate) feeding the dense query. The prompts and parsers are imported from Delphi’s source so they are identical by construction. What remains between the last rung and Delphi is exactly Delphi’s own contribution: exact-symbol, exact-path, path-affinity, and trigram branches; its chunker and index; tuned fusion weights; quoted-path demotion. Reported latencies exclude index construction for every system, as they did for Delphi.

Claim rule. “State of the art” would require leading every directly comparable engine on a C0 set with the complete comparable engine set runnable. No result in this paper meets it.

4 RQ2: REPOSITORY RETRIEVAL

Against lexical baselines. Delphi leads whole-file BM25, the official lexical ranker, and their tuned fusion on every set. On the independent final the Recall@20 leads over BM25 (+0.074 [+0.019, +0.139]) and the lexical ranker (+0.046 [+0.009, +0.083]) exclude zero while the fusion comparison is a tie; on SWE-bench MRR leads the strongest lexical system by +0.274 [+0.194, +0.412] with Recall@20 tied; on the C2 ARB partition every lexical comparison excludes zero. None of this is in dispute, and none of it isolates Delphi’s contribution.

Table 1: Independent commit-to-files final (C0): 18 cases, 6 repositories, scored once. Any-gold is the number of cases with a gold file in the top 20; latency is the median query time in seconds. Δ rows are Delphi minus the named system.

System	MRR	R@5	R@20	BCY@8k	any-gold	s
Delphi (frozen)	0.508	0.551	0.750	0.454	15/18	3.7
Dense (same embedding)	0.455	0.523	0.796	0.347	15/18	0.5
Dense+BM25 RRF	0.579	0.639	0.870	0.403	16/18	0.2
+ Delphi's rerankers	0.545	0.644	0.870	0.454	16/18	2.3
+ expansion	0.619	0.662	0.852	0.546	16/18	4.1
Lexical+BM25 RRF	0.473	0.500	0.759	0.287	15/18	1.6
Lexical ranker	0.438	0.403	0.704	0.259	14/18	1.1
BM25	0.370	0.458	0.676	0.259	14/18	1.2
<i>Paired Delphi—system deltas, repository-cluster bootstrap 95% intervals; * interval excludes zero</i>						
Δ vs. Dense (same embedding)	+0.053 [-0.084, +0.190]	+0.028 [-0.148, +0.194]	-0.046 [-0.139, +0.000]	+0.106 [-0.074, +0.259]	15 vs 15	$p=1$
Δ vs. Dense+BM25 RRF	-0.070 [-0.222, +0.086]	-0.088 [-0.236, +0.056]	-0.120* [-0.231, -0.019]	+0.051 [-0.120, +0.222]	15 vs 16	$p=1$
Δ vs. + Delphi's rerankers	-0.037 [-0.123, +0.049]	-0.093 [-0.324, +0.083]	-0.120* [-0.231, -0.019]	+0.000 [-0.139, +0.194]	15 vs 16	$p=1$
Δ vs. + expansion	-0.111* [-0.200, -0.034]	-0.111 [-0.324, +0.046]	-0.102 [-0.213, +0.000]	-0.093 [-0.278, +0.074]	15 vs 16	$p=1$
Δ vs. Lexical+BM25 RRF	+0.035 [-0.211, +0.259]	+0.051 [-0.088, +0.194]	-0.009 [-0.148, +0.083]	+0.167 [-0.074, +0.472]	15 vs 15	$p=1$
Δ vs. Lexical ranker	+0.070 [-0.177, +0.276]	+0.148 [-0.046, +0.343]	+0.046* [+0.009, +0.083]	+0.194 [-0.046, +0.481]	15 vs 14	$p=1$
Δ vs. BM25	+0.138 [-0.067, +0.301]	+0.093 [-0.037, +0.278]	+0.074* [+0.019, +0.139]	+0.194 [-0.046, +0.481]	15 vs 14	$p=1$

Table 2: SWE-bench Verified localization (C0): 62 instances, 12 repositories, scored once. Columns as in Table 1.

System	MRR	R@5	R@20	BCY@8k	any-gold	s
Delphi (frozen)	0.680	0.704	0.737	0.638	48/62	3.6
Dense (same embedding)	0.254	0.324	0.591	0.210	39/62	0.5
Dense+BM25 RRF	0.296	0.452	0.648	0.234	43/62	2.0
+ Delphi's rerankers	0.575	0.598	0.616	0.533	41/62	3.7
+ expansion	0.597	0.623	0.657	0.573	44/62	6.2
Lexical+BM25 RRF	0.354	0.500	0.719	0.306	47/62	13.2
Lexical ranker	0.406	0.509	0.702	0.387	46/62	11.0
BM25	0.149	0.230	0.416	0.145	29/62	10.7
<i>Paired Delphi—system deltas, repository-cluster bootstrap 95% intervals; * interval excludes zero</i>						
Δ vs. Dense (same embedding)	+0.426* [+0.328, +0.492]	+0.380* [+0.262, +0.583]	+0.146 [+0.000, +0.288]	+0.427* [+0.340, +0.500]	48 vs 39	$p=0.064$
Δ vs. Dense+BM25 RRF	+0.384* [+0.275, +0.464]	+0.252* [+0.152, +0.442]	+0.089 [-0.038, +0.182]	+0.404* [+0.288, +0.517]	48 vs 43	$p=0.3$
Δ vs. + Delphi's rerankers	+0.104 [-0.065, +0.210]	+0.106 [+0.000, +0.263]	+0.121 [-0.033, +0.259]	+0.105 [+0.000, +0.190]	48 vs 41	$p=0.14$
Δ vs. + expansion	+0.083 [-0.107, +0.203]	+0.081 [-0.056, +0.251]	+0.080 [-0.105, +0.217]	+0.065 [-0.062, +0.171]	48 vs 44	$p=0.42$
Δ vs. Lexical+BM25 RRF	+0.325* [+0.216, +0.453]	+0.204* [+0.099, +0.455]	+0.018 [-0.073, +0.075]	+0.331* [+0.211, +0.552]	48 vs 47	$p=1$
Δ vs. Lexical ranker	+0.274* [+0.194, +0.412]	+0.195* [+0.107, +0.403]	+0.035 [-0.042, +0.111]	+0.251* [+0.157, +0.433]	48 vs 46	$p=0.75$
Δ vs. BM25	+0.530* [+0.450, +0.662]	+0.474* [+0.352, +0.726]	+0.321* [+0.189, +0.572]	+0.493* [+0.395, +0.669]	48 vs 29	$p=0.00055$

Against the matched ladder. Table 1 is the first of the two clean tests. The dense rung alone, using nothing but the embedding model Delphi already calls, ties Delphi on any-gold and exceeds it on Recall@20 (0.796 vs. 0.750). Adding BM25 by untuned reciprocal-rank fusion yields MRR 0.579 and Recall@20 0.870 against Delphi's 0.508 and 0.750; the Recall@20 gap ($-0.120 [-0.231, -0.019]$) excludes zero. Attaching Delphi's own rerankers and expansion reaches MRR 0.619 and BCY 0.546, and Delphi's MRR deficit against that rung ($-0.111 [-0.200, -0.034]$) also excludes zero. Six repository clusters is a small basis for an interval and we report the direction with that caveat; but the point is not that the conventional stack wins, it is that nothing Delphi adds on top of it is visible here.

Table 2 is the second clean test. Its point estimates lean the other way and its conclusion is the same. On issue text the dense rung alone is weak (MRR 0.254) and untuned fusion with BM25 helps little (0.296); Delphi's leads over both have intervals excluding zero. Attaching Delphi's rerankers to that hybrid moves MRR from 0.296 to 0.575, and expansion to 0.597: the reranking head, not the candidate branches, is where the margin over lexical systems comes from. Against the final rung Delphi leads by +0.083 MRR $[-0.107, +0.203]$, +0.080 Recall@20 $[-0.105, +0.217]$, and +0.065 BCY $[-0.062, +0.171]$, with any-gold 48 vs. 44 ($p=0.42$)—a tie. Across the two C0 sets the sign of Delphi's difference from the conventional stack flips, neither set resolves it in Delphi's favor, and the only resolved differences (independent Recall@20 and MRR) favor the conventional stack. On the C2 ARB partition (Table 3) Delphi's margin over the final rung is +0.103 MRR $[-0.008, +0.189]$ and +0.076 BCY $[-0.019, +0.157]$, with Recall@5 and Recall@20 tied and any-gold 163 vs. 148 ($p=0.08$); against every earlier rung the leads exclude zero. That is the set on which Delphi's components were selected, so a margin there is the expected direction of selection, not evidence. Its workflow breakdown is still informative as a hypothesis: Delphi's lead is concentrated in trace2code (MRR 0.579 vs. 0.149; Recall@20 0.895 vs. 0.368) and edit2ripple, whose queries carry file paths and symbol names that its exact-path, path-affinity, and symbol branches match directly, while on code2test and comment2context the conventional rung is equal or

Table 3: ARB round-3 partition (C2, validation evidence only): 220 cases, 19 repositories. Columns as in Table 1. This table is reported for completeness; no confirmatory claim rests on it.

System	MRR	R@5	R@20	BCY@8k	any-gold	s
Delphi (frozen)	0.332	0.476	0.648	0.296	163/220	4.3
Dense (same embedding)	0.168	0.218	0.458	0.118	115/220	0.5
Dense+BM25 RRF	0.195	0.247	0.533	0.141	137/220	0.5
+ Delphi’s rerankers	0.214	0.317	0.526	0.176	135/220	2.1
+ expansion	0.229	0.389	0.595	0.220	148/220	4.2
Lexical+BM25 RRF	0.160	0.204	0.518	0.124	131/220	1.9
Lexical ranker	0.144	0.215	0.478	0.138	123/220	1.4
BM25	0.128	0.177	0.425	0.075	113/220	1.3
<i>Paired Delphi—system deltas, repository-cluster bootstrap 95% intervals; * interval excludes zero</i>						
Δ vs. Dense (same embedding)	+0.164* [+0.049, +0.257]	+0.258* [+0.090, +0.401]	+0.190* [+0.029, +0.321]	+0.178* [+0.058, +0.286]	163 vs 115	$p=5.9e-08$
Δ vs. Dense+BM25 RRF	+0.137* [+0.036, +0.215]	+0.229* [+0.074, +0.366]	+0.115 [-0.039, +0.237]	+0.155* [+0.060, +0.246]	163 vs 137	$p=0.0022$
Δ vs. + Delphi’s rerankers	+0.118* [+0.009, +0.198]	+0.159* [+0.007, +0.286]	+0.122 [-0.031, +0.240]	+0.120* [+0.039, +0.203]	163 vs 135	$p=0.0011$
Δ vs. + expansion	+0.103 [-0.008, +0.189]	+0.087 [-0.100, +0.238]	+0.053 [-0.117, +0.195]	+0.076 [-0.019, +0.157]	163 vs 148	$p=0.082$
Δ vs. Lexical+BM25 RRF	+0.172* [+0.083, +0.235]	+0.272* [+0.148, +0.372]	+0.130* [+0.024, +0.200]	+0.172* [+0.098, +0.263]	163 vs 131	$p=0.00017$
Δ vs. Lexical ranker	+0.188* [+0.106, +0.246]	+0.261* [+0.140, +0.361]	+0.170* [+0.088, +0.226]	+0.158* [+0.076, +0.255]	163 vs 123	$p=2.4e-06$
Δ vs. BM25	+0.204* [+0.111, +0.277]	+0.299* [+0.157, +0.433]	+0.223* [+0.111, +0.310]	+0.221* [+0.129, +0.303]	163 vs 113	$p=5e-09$

ahead. An exploratory split of the 62 SWE-bench issues by whether they contain a traceback or a Python path does not reproduce that pattern (Delphi’s largest advantage is on issues with neither, $n=35$), so we record it as a hypothesis for a designed test, not a finding.

The 100-instance expansion. To widen the C0 basis beyond 62 SWE-bench instances we provisioned the frozen stack on 100 further Verified instances drawn under a fresh salt (Appendix B) and will score them once, with the ladder and lexical comparators over the same snapshots. *[pending: expansion results; if not complete at submission, the draw and protocol are fixed and the run is reported as in progress].*

Development, briefly. The development partitions drove preregistered ablations whose outcomes we keep regardless of direction: exact vector scan replaced HNSW (+0.020 Recall@20 [0.000, 0.063]); query expansion stayed on; a Gemini embedding swap was rejected (MRR -0.054 [-0.102 , -0.011]); and a file-level Okapi BM25 branch was rejected when its single approved configuration failed preregistered gates (MRR interval floor -0.050 , BCY floor -0.037 , median latency ratio 1.34) despite losing zero any-gold cases. It ships default-off. Two development findings moved code: review-comment queries retrieved the file the comment already quoted, so paths present in the query envelope are demoted after the cross-encoder; and a global `*.pb.go` exclusion had silently dropped a gold file, so agent-mode indexing now retains generated source under the existing size and binary guards. Both were selected on development or C2 evidence, which is exactly why they cannot be credited from Table 3.

5 RQ3: DETERMINISM, LIKE FOR LIKE

The pipeline was the variance. Repeating eight ARB queries ten times against the July stack gave a mean exact top-10 rate of 30.6%. Paired traces isolated uncached, unseeded LLM stages; insertion-order ties in fusion SQL; and approximate HNSW under concurrency. Total SQL ordering, single-flight sharing of concurrent identical hosted calls, an explicit seed, and a persistent stage cache take exact repeats to 100% in-process across two concurrent 8×10 packets and two 75-case repeats; a genuine cold restart still moved 7/8 lists until the cache was warm, after which 8/8 survive two restarts. We claim restart-durable repeatability from the first persisted answer, not that two clean installations agree on their first answer. Embedding APIs were not the problem: the small OpenAI model jitters at the fifth decimal and never changed top-ten membership across $N=20$ repeats; Gemini was bitwise deterministic.

Same quantity, every engine. Our earlier headline—98.0% byte-identical context for Delphi against 35.6% and 0.0% for the hosted engines—compared retrieval stability for two engines with generated-text stability for the third. Table 4 separates the layers. On retrieved-set stability the synthesis engine’s cited sources are identical across every pair (1.000), Delphi’s retrieved items across 98.0%, and the documentation engine’s across 54% (quality-guided) or 50% (oracle IDs). What differs for the synthesis engine is the layer above retrieval: its answer text never repeats byte-for-byte, and that variation flips the identifier-hit outcome in 7.1% of pairs, against 0% for Delphi. The defen-

Table 4: Like-for-like repeatability on ten documentation cases $\times$ ten runs (450 run pairs per engine). Retrieved-set columns use each engine’s observable retrieval unit: retrieved items for Delphi and the documentation engine, the set of cited URLs for the synthesis engine. Byte-exact is the delivered context text; hit agreement is whether the identifier-hit outcome agreed across the pair.

Engine	Set exact	Order exact	Set Jaccard	Byte exact	Hit agreement
Delphi (raw context, $k=5$)	0.980	0.980	0.993	0.980	1.000
Documentation engine (quality-guided IDs)	0.542	0.362	0.835	0.356	0.918
Documentation engine (oracle IDs)	0.496	0.258	0.806	0.253	0.909
Synthesis engine (cited-URL set)	1.000	1.000	1.000	0.000	0.929

Table 5: Matched-output-contract documentation comparison, 40 development cases (C1) $\times$ 3 repeats through one frozen synthesis stage. Deltas are case-cluster bootstrap 95% intervals against the no-retrieval control. The synthesis engine’s native single pass is shown for reference only.

Arm	Identifier hit	Per-repeat	Δ vs. control
Documentation engine retrieval + frozen synthesis	0.625	0.650/0.600/0.625	+0.133 [+0.008, +0.258]
Delphi retrieval + frozen synthesis	0.617	0.600/0.625/0.625	+0.125 [+0.000, +0.250]
Synthesis engine retrieval + frozen synthesis	0.583	0.625/0.525/0.600	+0.092 [−0.008, +0.200]
Synthesis engine, native synthesis (recorded)	0.575	single pass	—
Model alone (no retrieval)	0.492	0.475/0.475/0.525	—

sible statements are narrower than before: Delphi’s caching and total ordering make its retrieved text repeatable, and repeatable retrieved text makes downstream decisions repeatable; the documentation engine’s retrieval drifts; the synthesis engine’s retrieval does not drift, its prose does. We have not tested whether a caching wrapper would give the hosted engines the same repeatability, and expect that it would.

6 DOCUMENTATION: MATCHED CONTRACTS, ALL FOUR ARMS

Documentation retrieval is where hosted engines are strongest and where output contracts differ most. The synthesis engine returns a fluent answer with citations; the documentation engine and Delphi return raw context. Scored naively, the synthesis engine leads (identifier hit 0.575 against Delphi’s 0.400 and the documentation engine’s 0.375). That comparison conflates retrieval with the synthesis model’s parametric knowledge.

Matching the contract, properly. Our first attempt routed the two raw-context engines through a frozen answer-style synthesis stage (gpt-5.4-mini, 1,600 output tokens, one system prompt) and compared them with the synthesis engine’s *native* pass, which a reviewer correctly identified as only partially matched: the output format was aligned, the synthesis model and prompt were not. The engine’s raw-retrieval mode returned no sources during the recorded round, so we could not route its retrieval at the time. Its synthesized responses, however, carry the five retrieved source chunks it grounded on. We reconstructed that context from the recorded responses (five sources per case, mean 2,754 tokens), packed it under the same 8,000-token budget with the same routine as every other arm, and routed it through the same frozen stage, three repeats, no new hosted calls. Table 5 is now a four-arm comparison under one synthesis model, one prompt, one budget.

Three findings. First, on raw retrieved context the synthesis engine is ahead of Delphi: its reconstructed five-source context hits the gold identifier on 0.475 of cases against Delphi’s 0.400 ($k=5$) and 0.425 ($k=20$). Second, under the matched contract the four arms finish within 0.05 of one another and every engine-versus-engine interval includes zero (Table 6); the synthesis engine’s matched arm (0.583) is within 0.008 of its native pass, so its own synthesis added nothing our frozen stage did not. Third, the metric saturates: the bare model scores 0.492, so roughly 80% of every arm’s score is parametric knowledge, and the retrieval-over-floor deltas are +0.13, +0.13, and +0.09. We do not detect a statistically resolved difference between any two engines under this evaluation. That is the supported sentence; we withdraw “parity or better,” which an interval of

Table 6: Engine-versus-engine paired differences under the matched contract (case-cluster bootstrap 95% intervals; wins/losses/ties over 40 cases). Every interval includes zero.

Comparison	Δ	95% CI	W/L/T
Delphi – synthesis engine (matched)	+0.033	[−0.092, +0.158]	7/5/28
Delphi – synthesis engine (native)	+0.042	[−0.092, +0.183]	6/6/28
Delphi – documentation engine	−0.008	[−0.117, +0.092]	6/5/29
Documentation engine – synthesis engine (matched)	+0.042	[−0.083, +0.167]	10/7/23
Synthesis engine matched – native	+0.008	[−0.050, +0.075]	5/4/31

[−0.09, +0.16] does not establish. Establishing non-inferiority would require a predeclared margin and a sample several times this size.

7 RQ1: DOES ANY OF IT HELP AN AGENT?

Closed-tool seeding (a replication). ARB’s own seed intervention holds a closed-tool agent fixed and varies its initial context. We replicated it with the policy model `gpt-5.4-mini`, tools `list_dir/grep/read_file/submit`, and 40 paired development cases per arm, adding a July-Delphi seed. File-F1: no seed 0.046; random files 0.013; BM25 seed 0.093; lexical seed 0.119; July-Delphi seed 0.234; oracle 0.858. The Delphi seed beats no-seed by +0.188 [0.055, 0.294] and BM25 by +0.142 [0.059, 0.208]; against lexical the F1 interval includes zero [−0.002, 0.190]. Seeds change which files a frozen agent submits, ordered by seed quality, which is what the benchmark’s authors also found. The oracle ceiling shows localization headroom in this harness; it does not show that retrieval dominates reasoning, editing, or validation failures in a repairing agent, and we withdraw that sentence from the earlier draft.

DS-1000 (unresolved, not negative). With the frozen generator, 40 cases $\times$ 3 repeats: no retrieval 0.575; documentation engine 0.567 (−0.008 [−0.142, 0.125]); Delphi retrieval+synthesis 0.592 (+0.017 [−0.125, 0.158]); synthesis engine 0.642 (+0.067 [−0.075, 0.200]). Every interval spans effects from substantial harm to substantial benefit. This experiment does not show that retrieval fails to help; it shows that 40 cases cannot tell.

Executable pilot (preregistered). The reviewers’ request was direct: measure retrieval and executable outcome on the same tasks. We preregistered a pilot (Appendix C) on the 62 C0 SWE-bench instances whose rankings every system in Table 2 had already produced once. The agent is mini-SWE-agent with `gpt-5.4-mini`, ordinary shell tools in every condition, and a fixed step budget; a seed is a block appended to the issue naming the top five files a recorded ranking produced, with the head of each file, capped at 3,000 tokens. Conditions: no seed; five random non-gold files in the same block (prompting control); Delphi’s recorded top five; the conventional ladder’s recorded top five. The outcome is verified repair under the official harness, paired by instance, with instance- and repository-cluster intervals and exact McNemar tests. Table 7 reports the first repeat at step budget 50; a second repeat of every condition is running and will be reported alongside it.

Four observations, all provisional at one trajectory per cell. First, a Delphi seed raised verified repairs from 25 to 31 of 62 (+0.097 [+0.016, +0.194]; McNemar 7 vs. 1, $p=0.07$). Second, the random-file control also rose, to 30 (+0.081 [−0.016, +0.177]), and Delphi’s difference from it is +0.016 [−0.065, +0.097]: at this size we cannot distinguish the value of *which* files were named from the value of the block itself, which is exactly the confound the control exists to expose. Third, the conventional ladder’s seed, whose top five named a gold file on 42 instances against Delphi’s 47, resolved only 23 (−0.129 vs. Delphi [−0.242, −0.032]; 10 vs. 2, $p=0.04$), and fewer than the random control; when its seed named a gold file the agent repaired 17/42 against 27/47 for Delphi. We do not have an explanation and will not offer one until the second repeat says whether the gap is stable or sampling noise. Fourth, budgets were not binding: agents used 7–9 steps of 50 and cost 1.6–2.6 cents per instance, so the seed’s own tokens raised cost by 30–60% while Delphi’s seed reduced steps relative to the random block (−1.7 [−3.8, −0.2]). The pilot’s width is the finding as much as its point estimates: at $n=62$ and one trajectory, effects of ± 0.08 are inside the noise, so the

Table 7: Executable pilot, first repeat: 62 C0 SWE-bench Verified instances, mini-SWE-agent with gpt-5.4-mini, step budget 50, one trajectory per instance and condition. Resolved is the official harness verdict. * instance-cluster interval excludes zero.

Condition	Resolved (of 62)	Rate	Mean cost (USD)	Mean steps
No seed	25	0.403	0.0162	7.7
Random files (control)	30	0.484	0.0263	9.0
Delphi top-5 (frozen)	31	0.500	0.0208	7.3
Conventional ladder top-5	23	0.371	0.0240	7.5
<i>Paired differences in resolved rate: instance-cluster 95% CI [repository-cluster CI]; discordant a/b; exact McNemar p</i>				
Delphi – no seed	+0.097*	[+0.016, +0.194]	[+0.019, +0.244]; 7/1; $p=0.07$	
Delphi – random	+0.016	[−0.065, +0.097]	[−0.057, +0.154]; 4/3; $p=1.00$	
Delphi – conventional	+0.129*	[+0.032, +0.242]	[+0.029, +0.317]; 10/2; $p=0.04$	
Conventional – no seed	−0.032	[−0.113, +0.048]	[−0.147, +0.045]; 2/4; $p=0.69$	
Conventional – random	−0.113*	[−0.210, −0.032]	[−0.208, −0.059]; 1/8; $p=0.04$	
Random – no seed	+0.081	[−0.016, +0.177]	[−0.015, +0.200]; 7/2; $p=0.18$	

study that could answer the title question needs several hundred instances and repeated trajectories, which the trivially low per-run cost makes affordable.

8 THE HOSTED REPOSITORY ARM

The synthesis engine also markets a repository surface. Scoring it requires 68 exact (repo, base_commit) snapshots in the provider’s account for the development split alone. Whole-snapshot ingestion of large repositories stalled in a non-terminal syncing state; 32/68 pairs remain indexed from an earlier sharded round; a minimal one-shard probe failed to leave syncing across a 40-minute deadline and a three-hour watch in August, and a fresh one-shard probe issued on 2026-09-09 with a newly provided credential (which resolves to the same account) timed out in the same state after 40 minutes; the account inventory still shows the August probes syncing. No repository score is reported for that engine in either direction. This is a gap in the evidence, not a result, and after the matched-ladder findings it is no longer the most important gap.

9 WHAT WE CLAIM, AND WHAT WE DO NOT

Claimed. (1) Delphi leads whole-file BM25 and the official lexical ranker on every set, with intervals excluding zero on the C0 sets for Recall@20 (independent) and MRR (SWE-bench). (2) A conventional dense+BM25 hybrid built from Delphi’s own embedding model, with Delphi’s own rerankers and expansion attached, matches or exceeds Delphi on the C0 sets; on the independent final two of its leads have intervals excluding zero. (3) Under one frozen synthesis stage, no two documentation engines are statistically distinguishable at $n=40$, and all sit 0.08–0.13 above a no-retrieval floor. (4) Delphi’s retrieved text is repeatable in-process and across restarts with a warm cache; on retrieved-set stability the hosted synthesis engine is at least as stable. (5) Retrieval seeds change a closed-tool agent’s file selection, ordered by seed quality, replicating ARB.

Not claimed. (1) State of the art, anywhere. (2) Any confirmatory result on the ARB round-3 partition, which is C2. (3) Documentation parity or non-inferiority. (4) Downstream utility on DS-1000 at $n=40$, in either direction. (5) That Delphi’s structural branches, tuned weights, or quoted-path demotion add value over the conventional ladder: on the sets that can bear the question, they do not separate from it.

10 CONCLUSION

Measured against baselines built from its own parts, and on cases its development could not see, Delphi’s advantage reduces to the components any careful practitioner would assemble: an embedding model, BM25, reciprocal-rank fusion, a small cross-encoder, a listwise reranker, and query expansion. The claims we made from the re-partitioned ARB set did not survive the ledger that classified how that set had been used, and the claims we made against hosted engines did not survive

matching the quantity being measured. What survives is a protocol—exposure classes at the case level, component-matched baselines, matched output contracts, like-for-like determinism, artifact-generated tables—and one open question the pilot begins to answer: whether any of this changes what an agent repairs.

AI USE STATEMENT

Generative AI tools were used to help formulate experimental questions, write and test evaluation and product code, operate experiments, inspect failures, search for related work, and draft and edit this manuscript. They were not used to generate benchmark gold labels. Empirical claims are checked against persisted run artifacts and canonical repository snapshots; every table in this paper is generated from those artifacts by a script in the released archive. All AI-assisted work remains subject to human review before submission, and the authors take responsibility for the final content, claims, code, and artifacts.

ETHICS STATEMENT

The study uses public open-source repositories and public benchmark records. Repository contents are processed locally except where a named hosted retrieval or model API is itself the evaluated system. We retain commit identifiers and public issue or review text for provenance, but do not intentionally collect private repository data or credentials. Because benchmark results can be used in product marketing, we predeclare the claim rule, preserve failed runs, publish the exposure ledger, and give losses the same prominence as wins.

REPRODUCIBILITY STATEMENT

The archive accompanying this paper contains the split manifests and case-level exposure ledger, every per-case artifact (queries, gold paths, ranked outputs, synthesized answers, latencies), every paired-analysis JSON, the harness code for every system and comparator including the matched ladder, the frozen serving configuration and the scripts that launch it, the pilot protocol and trajectories, and the script that regenerates every table in this paper from those artifacts. Repository corpora are re-cloned from public GitHub by a provided script; hosted-engine runs require the reader’s own credentials.

REFERENCES

- Zhaoling Chen, Robert Tang, Gangda Deng, Fang Wu, Jialong Wu, Zhiwei Jiang, Viktor Prasanna, Arman Cohan, and Xingyao Wang. Locagent: Graph-guided llm agents for code localization. In *ACL*, pp. 8697–8727, 2025. doi: 10.18653/v1/2025.acl-long.426.
- Luyu Gao, Xueguang Ma, Jimmy Lin, and Jamie Callan. Precise zero-shot dense retrieval without relevance labels. In *ACL*, 2023.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *ICLR*, 2024.
- Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *EMNLP*, 2020.
- Omar Khattab and Matei Zaharia. ColBERT: Efficient and effective passage search via contextualized late interaction over BERT. In *SIGIR*, 2020.
- Han Li, Letian Zhu, Bohan Zhang, Rili Feng, Jiaming Wang, Yue Pan, Earl T. Barr, Federica Sarro, Zhaoyang Chu, and He Ye. Contextbench: A benchmark for context retrieval in coding agents. *arXiv preprint arXiv:2602.05892*, 2026.
- Bowen Qin and Yi Xie. Agent retrieval bench: Evaluating repository context retrieval for coding agents. *arXiv preprint arXiv:2607.24882*, 2026.

- Stephen Robertson and Hugo Zaragoza. The probabilistic relevance framework: Bm25 and beyond. *Foundations and Trends in Information Retrieval*, 3(4):333–389, 2009.
- Chunqiu Steven Xia, Yinlin Deng, Soren Dunn, and Lingming Zhang. Agentless: Demystifying llm-based software engineering agents. *arXiv preprint arXiv:2407.01489*, 2024.
- John Yang, Carlos E. Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. Swe-agent: Agent-computer interfaces enable automated software engineering. In *NeurIPS*, 2024.
- Fengji Zhang, Bei Chen, Yue Zhang, Jacky Keung, Jin Liu, Daoguang Zan, Yi Mao, Jian-Guang Lou, and Weizhu Chen. Repocoder: Repository-level code completion through iterative retrieval and generation. In *EMNLP*, pp. 2471–2484, 2023. doi: 10.18653/v1/2023.emnlp-main.151.
- Fuwei Zhang, Yanzhao Zhang, Mingxin Li, Dingkun Long, Lexiang Hu, Pengjun Xie, Zhao Zhang, and Fuzhen Zhuang. Core-bench: A comprehensive benchmark for code retrieval in the era of agentic coding. *arXiv preprint arXiv:2606.11864*, 2026a.
- Shaoqiu Zhang, Yuhang Wang, Jialiang Liang, Yuling Shi, Wenhao Zeng, Maoquan Wang, Shilin He, Ningyuan Xu, Siyu Ye, Kai Cai, and Xiaodong Gu. Swe-explore: Benchmarking how coding agents explore repositories. *arXiv preprint arXiv:2606.07297*, 2026b.
- Sheng Zhang, Yifan Ding, Shuquan Lian, Shun Song, and Hui Li. Coderag: Finding relevant and necessary knowledge for retrieval-augmented repository-level code completion. In *EMNLP*, pp. 23278–23288, 2025. doi: 10.18653/v1/2025.emnlp-main.1187.

A FROZEN SYSTEM SPECIFICATION

Frozen commit 91d76c1 of the Delphi repository (branch feature/generated-source-freeze), served natively against PostgreSQL with pgvector. Every scored response carried the following attested configuration: embedding model text-embedding-3-small; vector mode exact scan (HNSW disabled for scoring; ef_search 100 when enabled); hybrid candidate branches vector, BM25, exact symbol, exact path, path affinity, and trigram with reciprocal-rank fusion weights 0.50/0.25/0.20/0.15/0.15/0.05 and $k=60$; 50 fused candidates; cross-encoder cross-encoder/ms-marco-MiniLM-L-6-v2 over the top 30 with blend $\alpha=0.4$ (code-aware reranker disabled); quoted-path demotion when the query envelope carries evidence keys; listwise reranking of the top 20 by gpt-4o at temperature 0, seed 1042, 280-character excerpts, 300-token replies; hypothetical-document expansion by gpt-4o-mini (temperature 0, seed 1042, 220-token replies) gated on prose-like queries and feeding only the vector branch; persistent SQLite stage cache; API top- $k=20$; file-diverse BM25 and path-token branches disabled; agent quality mode with tests, docs, examples, configs, and manifests indexed and generated source retained under 500,000-byte and NUL-content guards. The listwise and expansion system prompts are the strings in synsc/services/listwise_rerank.py and synsc/services/query_expansion.py at the frozen commit; the matched ladder imports them from the same files. Context packing for BCY@8k uses the official ARB pack_files at an 8,000-token budget.

B EXPOSURE LEDGER

Set	Cases	Class	Supports
ARB round-3 partition (positive workflows)	220	C2	validation only
ARB round-3 development	75	C1	development
Independent commit-to-files final	18	C0	confirmatory
Independent commit-to-files development	48	C1	development
SWE-bench Verified, round 3	62	C0	confirmatory
SWE-bench Verified, round-4 expansion	100	C0	confirmatory (scoring in progress)
DS-1000 documentation subset	40	C1	development

Decisions and the evidence that informed them: query-time embedding alignment, RRF fusion, the path-affinity branch, the cross-encoder blend, expansion, and listwise reranking were accepted on round-2 development (75) and round-2 held-out (220) aggregates over the ARB v2 pool that round 3 re-partitioned; quoted-path demotion on July per-workflow rates over all 427 cases; exact scan, single-flight, the persistent cache, and total SQL ordering on round-3 development and determinism probes; generated-source indexing on the gold-path preflight over the independent and ARB development sets; the File-Okapi rejection on both development sets. The machine-readable ledger lists every case ID per set.

C EXECUTABLE PILOT PROTOCOL

Tasks: the 62 round-3 SWE-bench Verified instances. Agent: mini-SWE-agent 2.4.6, text-based action format, `gpt-5.4-mini`, official `x86_64` instance images. Conditions: `none`; `random` (five candidate files drawn with seed 1042 per instance, gold excluded); `delphi` (top five of the frozen stack’s recorded ranking); `hybrid_rerank_expand` (top five of the conventional ladder’s recorded ranking). Seed block: identical wording, up to five paths with the first 40 lines of each file at the base commit, capped at 3,000 tokens. Budgets: primary step limit 50 with a 2.0 USD cost cap; secondary step limit 15 if time allows. Outcome: verified repair under `swebench 5.0.2` against `SWE-bench/SWE-bench_Verified`; secondary outcomes API cost, steps, whether a gold file was opened, and submission rate. Analysis: paired by instance; instance-cluster and repository-cluster bootstrap 95% intervals (20,000 resamples, seed 1042); exact McNemar on discordant pairs. An interval including zero at $n=62$ is reported as unresolved.